

Model Training, Evaluation, and Results

This document outlines the steps involved in training, evaluating, and interpreting the results of a classification model using the UCI Adult Census dataset.

1. Splitting the Data into Training and Testing Sets

```
```python
X_train, X_test, y_train, y_test = train_test_split(X_train, y_train, test_size=0.2, random_state=42)
```
```

The dataset is split into training and testing sets using an 80-20 split. With `test_size=0.2` indicating that 20% of the data is reserved for testing. A fixed `random_state` ensures reproducibility of results. This step is crucial for evaluating the model's performance on unseen data, which helps in assessing its generalization capability and combat over or under fitting.

2. Training the Model Using the Most Important Features

```
```python
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)
```
```

A `RandomForestClassifier` is instantiated with 100 estimators and a fixed `random_state` for reproducibility. The model is then trained on the training set (`X_train`, `y_train`). The use of the most important features, identified during the preprocessing step, ensures that the model is efficient and focuses on the variables that contribute the most to the prediction task. Random forests was chosen for its robustness, ability to handle various types of data, and resistance to overfitting.

3. Making Predictions on the Testing Set

```
```python
y_pred = rf.predict(X_test)
```
```

Predictions are made on the testing set (`X_test`) using the trained model. This step evaluates the model's performance on unseen data, providing an indication of its generalization ability and giving us time to tune the model to address its shortcomings.

4. Calculating the Accuracy Score

```
```python
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
...
```

The accuracy score, which is the ratio of correctly predicted instances to the total instances, is calculated using `accuracy\_score` from `scikit-learn`. This metric provides a straightforward measure of the model's overall performance. It is to be noted that this alone will not give us a complete picture of the model performance and other metrics like F1-Score are to also be accounted for.

## 5. Generating the Classification Report

```
```python
```

```
report = classification_report(y_test, y_pred, output_dict=True)
```

```
```
```

A detailed classification report is generated using `classification\_report` from `scikit-learn`. The report includes precision, recall, and F1-score for each class, as well as macro and weighted averages of these metrics. These metrics provide a comprehensive evaluation of the model's performance, highlighting its ability to correctly predict each class and its overall balance. This is a much more exact report of the models performance and should be considered more deeply than the simple accuracy score.

### Output

**Accuracy: 0.846668791839337 = ~85%**

#### **Classification Report:**

##### **<=50K:**

precision: 0.8581652569941445

recall: 0.8337547408343868

f1-score: 0.8457839050977878

support: 1582.0

##### **>50K:**

precision: 0.835625,

recall: 0.8598070739549839,

f1-score: 0.8475435816164818,

support: 1555.0

##### **macro avg:**

precision: 0.8468951284970723

recall: 0.8467809073946854

f1-score: 0.8466637433571348

support: 3137.0

##### **weighted avg:**

precision: 0.8469921299218159

recall: 0.846668791839337

f1-score: 0.8466561706338316

support: 3137.0

## Interpretation:

**Accuracy:** The model achieves an accuracy of approximately 84.67%, indicating that it correctly predicts the income class for a significant portion of the test instances.

**Precision, Recall, and F1-score:** The classification report provides detailed metrics for both classes ['<=50K' and '>50K']. Precision measures the proportion of true positive predictions among all positive predictions, recall measures the proportion of true positive predictions among all actual positives, and F1-score is the harmonic mean of precision and recall. The metrics show balanced performance across both classes.

For the '<=50K' class, the model achieves a precision of approximately 85.82%, a recall of approximately 83.38%, and an F1-score of approximately 84.58%.

For the '>50K' class, the model achieves a precision of approximately 83.56%, a recall of approximately 85.98%, and an F1-score of approximately 84.75%.

Given the macro and weighted output values they indicate consistent performance throughout.